

Fraud Detection System - Take-Home Test




This project implements a complete machine learning solution for detecting fraudulent transactions across four main deliverables.

Table of Contents

- [Part 1: Exploratory Data Analysis \(EDA\)](#)
- [Part 2: Model Training](#)
- [Part 3: Model Serving \(API and Storage\)](#)
- [Part 4: System Architecture](#)
- [Setup Instructions](#)
- [Running the Complete Workflow](#)

Take-Home Test Deliverables

All Parts Completed:

1.  **Part 1: Exploratory Data Analysis (EDA)** - [Training.ipynb](#)
2.  **Part 2: Model Training** - [Training.ipynb](#)
3.  **Part 3: Model Serving (API and Storage)** - [app.py](#) + [streamlit_app.py](#)
4.  **Part 4: System Architecture Design** - Architecture diagram and explanation

Part 1: Exploratory Data Analysis (EDA)

Location: [Training.ipynb](#)

Deliverable: Comprehensive analysis of 6.36M transactions revealing key fraud patterns and feature engineering insights.

Part 2: Model Training

Location: [Training.ipynb](#)

Deliverable: XGBoost classifier achieving 98% precision and 99% recall on fraud detection.

Part 3: Model Serving (API and Storage)

Components:

- **Flask REST API** ([app.py](#)) with fraud prediction and storage
- **Streamlit Web Interface** ([streamlit_app.py](#)) for interactive testing
- **SQLite Database** for persistent fraud case storage

API Endpoints:

```
POST /predict      # Fraud prediction with automatic flagging
GET /frauds        # Paginated fraud history (page, per_page params)
GET /health        # API health check
```

Key Features:

- **Automatic Feature Engineering:** API calculates `error_bal_src/dst` internally
- **Fraud Storage:** Automatically stores flagged transactions with metadata
- **Pagination:** Efficient handling of large fraud case volumes
- **Web Testing Interface:** User-friendly testing with predefined examples

Test Examples:

```
// Fraudulent TRANSFER (empties source account)
{
  "time_ind": 1, "transac_type": "TRANSFER", "amount": 181.0,
  "src_bal": 181.0, "src_new_bal": 0.0,
  "dst_bal": 0.0, "dst_new_bal": 181.0
}
// Result: {"is_fraud": 1}

// Legitimate PAYMENT
{
  "time_ind": 1, "transac_type": "PAYMENT", "amount": 100.0,
  "src_bal": 1000.0, "src_new_bal": 900.0,
  "dst_bal": 0.0, "dst_new_bal": 0.0
}
// Result: {"is_fraud": 0}
```

Dependencies

System Requirements:

- **Docker:** Latest version with Docker Compose
- **Docker Compose:** v2.0 or higher
- **Available Ports:** 8888 (training), 5001 (API), 8501 (web interface)
- **Disk Space:** ~2GB for Docker images and dataset
- **Memory:** 4GB RAM recommended for training

All Python dependencies are automatically handled by Docker containers:

- **Core API:** Flask, scikit-learn, xgboost, pandas, joblib
- **Training:** Jupyter, numpy, matplotlib, seaborn, gdown
- **Web Interface:** Streamlit, requests

No manual Python installation required - everything runs in Docker!

Setup Instructions

Prerequisites: Ensure Docker and Docker Compose are installed on your system.

Step 1: Clone Repository

```
git clone https://github.com/tumrabort/FraudDetecion
cd FraudDetecion
```

That's it! No additional setup required - Docker handles everything.

Running the Complete Workflow

Run EDA and Model Training (For part 1&2)

Single Docker Command:

```
# Start Jupyter Lab environment for training
docker-compose -f docker-compose.training.yml up --build
```

Access the training environment:

- **URL:** `http://localhost:8888`
- **Token:** `fraud_detection_training`

Training Instructions (CRITICAL - Follow Exactly):

1. Wait for the Docker container to fully start (you'll see "Jupyter server is running")
2. Open your browser and go to `http://localhost:8888`
3. Enter token: `fraud_detection_training`
4. Click on `Training.ipynb` to open the notebook
5. **Execute ALL cells in sequence** - this will:
 - Download 6.36M transaction dataset from Google Drive
 - Perform comprehensive EDA revealing key fraud patterns
 - Engineer features and handle class imbalance
 - Train XGBoost model with 98% precision and 100% recall
 - Save the complete model pipeline to `model/fraud_model.joblib`
6. **IMPORTANT:** Ensure the final cell saves the model to `model/fraud_model.joblib`
7. When training is complete, stop the container: `Ctrl+C`

Expected Output: A trained model file at `model/fraud_model.joblib` (271KB)

Step 3: Model Serving (API and Storage)

Prerequisites: Ensure Step 1-2 is completed and `model/fraud_model.joblib` exists.

Single Docker Command:

```
# Start the complete API stack with web interface
docker-compose -f docker-compose.api.yml up --build
```

Expected Startup Output:

```
fraud-detection-api-1 | Model loaded successfully.
fraud-detection-api-1 | * Running on all addresses (0.0.0.0)
fraud-detection-api-1 | * Running on http://127.0.0.1:5000
streamlit-ui-1         | You can now view your Streamlit app in your
streamlit-ui-1         | browser.
streamlit-ui-1         | URL: http://0.0.0.0:8501
```

Access the services:

- **Fraud Detection API:** <http://localhost:5001>
- **Streamlit Web Interface:** <http://localhost:8501>

Services included:

- Flask REST API with trained XGBoost model
- SQLite database for persistent fraud storage
- Streamlit web interface for interactive testing
- Automatic health checks and restart policies
- Automatic health checks and restart policies

Testing the System


Method 1: Streamlit Web Interface (Recommended)

Step-by-step testing procedure:


1. Verify services are running:

- API should be accessible at <http://localhost:5001>
- Web interface should be accessible at <http://localhost:8501>

2. Open the web interface:

- Navigate to <http://localhost:8501> in your browser
- You should see "  Fraud Detection API Tester"

3. Test API health:

- Click "Check API Health" button
- Expected result:  API is healthy with `{"model_loaded": true}`

4. Test fraud prediction:

- Use the "Predefined Examples" tab

- Select "Fraudulent Transfer" and click "📝 Test Selected Example"
- Expected result: 🚨 **FRAUD DETECTED**
- Select "Legitimate Payment" and test
- Expected result: ✅ **LEGITIMATE**

5. View flagged transactions:

- Click "🔄 Refresh Flagged Transactions"
- Should show previously flagged fraudulent transactions

Method 2: Direct API Testing (Advanced)

Health Check:

```
curl http://localhost:5001/health
# Expected: {"status": "healthy", "model_loaded": true}
```

Test Fraud Detection:

```
# Test 1: Fraudulent TRANSFER (empties source account)
curl -X POST -H "Content-Type: application/json" -d '{
  "time_ind": 1,
  "transac_type": "TRANSFER",
  "amount": 181.0,
  "src_bal": 181.0,
  "src_new_bal": 0.0,
  "dst_bal": 0.0,
  "dst_new_bal": 181.0
}' http://localhost:5001/predict
# Expected: {"is_fraud": 1}

# Test 2: Legitimate PAYMENT
curl -X POST -H "Content-Type: application/json" -d '{
  "time_ind": 1,
  "transac_type": "PAYMENT",
  "amount": 100.0,
  "src_bal": 1000.0,
  "src_new_bal": 900.0,
  "dst_bal": 0.0,
  "dst_new_bal": 0.0
}' http://localhost:5001/predict
# Expected: {"is_fraud": 0}
```

View Stored Fraudulent Transactions:

```
curl "http://localhost:5001/frauds?page=1&per_page=10"
# Expected: JSON with fraudulent_transactions array and pagination info
```

Complete Docker Deployment Guide

Training Environment

```
# Start training environment
docker-compose -f docker-compose.training.yml up --build

# Access at: http://localhost:8888 (Token: fraud_detection_training)
# Complete the training, then stop:
docker-compose -f docker-compose.training.yml down
```

Production API Stack

```
# Start complete API stack
docker-compose -f docker-compose.api.yml up --build

# Access services:
# - API: http://localhost:5001
# - Web UI: http://localhost:8501
# Stop when done:
docker-compose -f docker-compose.api.yml down
```

Troubleshooting

If API fails to start:

1. Ensure training step completed successfully
2. Verify `model/fraud_model.joblib` exists (should be ~271KB)
3. Check Docker logs: `docker-compose -f docker-compose.api.yml logs`

If ports are in use:

```
# Check what's using the ports
lsof -i :8888 -i :5001 -i :8501
# Kill processes or change ports in docker-compose files
```

Clean restart:

```
# Remove all containers and rebuild
docker-compose -f docker-compose.training.yml down
docker-compose -f docker-compose.api.yml down
docker system prune -f
# Then restart the services
```

Part 4: System Architecture

Overview

This section designs a **production-ready fraud detection system** that processes real-world banking transactions at scale. Unlike the current demo API that handles individual requests, a production system must:

- **Process millions of transactions per day** in real-time
- **Integrate with existing banking infrastructure** (core banking systems, payment processors)
- **Provide tools for human auditors** to investigate flagged cases
- **Maintain high availability** with zero tolerance for downtime
- **Scale automatically** during peak transaction periods (e.g., Black Friday, payroll days)

Real-World Context: Banks generate 10,000-50,000+ transactions per second during peak hours. Each transaction must be evaluated for fraud **before completion** (typically within 50-100ms), making this a high-throughput, low-latency challenge.

System Architecture

```
graph TB
    subgraph "External Systems"
        BANK[Banking Systems  
ATM, Mobile, Web]
        KAFKA[Kafka Cluster  
transactions topic  
3 brokers, 6 partitions]
    end

    subgraph "Fraud Detection Layer"
        CG[Consumer Group  
6 instances  
At-least-once delivery]
        subgraph "Processor"
            PROCESSOR[Transaction Processor  
Feature Engineering]
        end
        MODEL[ML Model Service  
XGBoost Pipeline  
In-memory]
    end

    subgraph "Data Storage"
        CACHE[(Redis Cache  
Recent predictions)]
        DB[(PostgreSQL  
Fraud cases  
Read replicas)]
    end

    subgraph "Auditor Interface"
        LB[Load Balancer]
```

```

NGINX]
    API[REST API
Case Management]
    WEB[Web Dashboard
Investigation UI]
end

    subgraph "Monitoring"
    METRICS[Prometheus
Grafana]
    ALERTS[Alerting
Slack/Email]
end

BANK --> KAFKA
KAFKA --> CG
CG --> PROCESSOR
PROCESSOR --> MODEL
MODEL --> CACHE
MODEL --> DB

LB --> API
API --> DB
API --> CACHE
WEB --> LB

MODEL --> METRICS
API --> METRICS
METRICS --> ALERTS

```

Component Details

1. Kafka Integration - Transaction Streaming

Why Kafka?

Banking systems need **event-driven architecture** to handle transaction volumes. Kafka provides:

- **Guaranteed message delivery** (critical for financial data)
- **Horizontal scaling** through partitioning
- **Replay capability** for reprocessing during system failures
- **Decoupling** between transaction producers (ATMs, mobile apps) and fraud detection

Consumer Strategy Explained:

- **6 Consumer Instances:** Match Kafka's 6 partitions for optimal throughput
- **Partition by Account ID:** Ensures all transactions for an account are processed in order (prevents race conditions)
- **At-least-once Delivery:** Critical for fraud detection - better to check twice than miss fraud
- **Dead Letter Queue:** Failed messages don't block the pipeline; they're investigated separately

Message Format:


```
{
  "transaction_id": "TXN_001",
  "timestamp": "2024-01-15T10:30:00Z",
  "type": "TRANSFER",
  "amount": 50000.00,
  "src_account": "ACC_123",
  "dst_account": "ACC_456",
  "src_balance": 100000.00,    // Balance before transaction
  "dst_balance": 25000.00    // Balance before transaction
}
```

2. Fraud Detection Service - The Brain of the System

Why This Architecture?

Traditional rule-based systems (e.g., "flag transactions >\$10,000") generate **80%+ false positives**. Our ML-based approach provides:

- **Context-aware detection:** Considers account history, transaction patterns
- **Adaptive learning:** Model improves as fraud patterns evolve
- **Precision targeting:** 98% precision means minimal disruption to legitimate customers

How the Processing Pipeline Works:

```
graph LR
    A[Kafka Message  
Raw Data  
] --> B[Validation  
Clean Data  
]
    B --> C[Feature Engineering  
ML Features  
]
    C --> D[ML Model  
0.85 Score  
]
    D --> E[Decision  
FRAUD  
]
    E --> F[Storage  
Database  
]

    style A fill:#e1f5fe,color:#000000
    style B fill:#f3e5f5,color:#000000
    style C fill:#e8f5e8,color:#000000
    style D fill:#fff3e0,color:#000000
    style E fill:#ffebee,color:#000000
    style F fill:#f1f8e9,color:#000000
```

Step-by-Step Breakdown:

1. Message Validation (10ms):

- **Schema Validation:** Ensure all required fields present
- **Data Quality:** Check for impossible values (negative balances, future timestamps)
- **Business Rules:** Verify account exists, transaction type valid
- **Why Critical:** Garbage in = garbage out; prevents model degradation

2. Feature Engineering (15ms):

- **Real-time Calculations:** $\text{error_bal_src} = \text{src_bal} - \text{amount} - \text{src_new_bal}$ and $\text{error_bal_dst} = \text{dst_bal} + \text{amount} - \text{dst_new_bal}$
- **Historical Context:** Account velocity, typical transaction amounts
- **Pattern Recognition:** Time-of-day patterns, geographic anomalies
- **Why In Real-Time:** Features must reflect current transaction state

3. Model Inference (20ms):

- **XGBoost Prediction:** Our trained model processes engineered features
- **Confidence Scoring:** 0.0-1.0 probability score (0.85 = 85% fraud probability)
- **Ensemble Logic:** Could combine multiple models for robustness
- **Memory Optimization:** Model loaded in RAM for sub-millisecond inference

4. Decision Logic (5ms):

- **Threshold Application:** >0.5 probability = fraud flag
- **Risk-Based Actions:** High-risk (>0.9) = immediate block, medium-risk = additional auth
- **Business Rules Override:** VIP customers, trusted merchants get different thresholds

5. Storage & Alerting (10ms):

- **Audit Trail:** Every decision logged for regulatory compliance
- **Real-time Alerts:** High-risk cases trigger immediate investigator notification
- **Case Creation:** Automatic assignment to fraud analysts

Performance Requirements Explained:

- **<50ms Latency (p95):** Customer expects instant transaction approval
- **10,000+ TPS:** Peak banking hours (lunch time, end-of-month payrolls)
- **>95% Precision:** False positives cost customer satisfaction + ops overhead
- **>98% Recall:** Missing fraud costs money + regulatory penalties

Horizontal Scaling Strategy:

```
graph TD
  LB[Load Balancer] --> FS1[Fraud Service 1]
  LB --> FS2[Fraud Service 2]
  LB --> FS3[Fraud Service 3]
  LB --> FSN[... Auto-scale 2-15 instances  
based on Kafka lag]
```

```

FS1 --> DB[(Database)]
FS2 --> DB
FS3 --> DB
FSN --> DB

style LB fill:#e3f2fd,color:#000000
style FS1 fill:#f3e5f5,color:#000000
style FS2 fill:#f3e5f5,color:#000000
style FS3 fill:#f3e5f5,color:#000000
style FSN fill:#f3e5f5,color:#000000
style DB fill:#e8f5e8,color:#000000

```

3. Database Design - Handling Millions of Fraud Cases

Why This Database Strategy?

Fraud detection generates **massive data volumes** (5-10% of transactions flagged = 50,000+ cases/day).

The system needs:

- **Fast writes** for real-time case creation
- **Complex queries** for auditor investigations
- **Historical analysis** for pattern detection
- **Regulatory compliance** with audit trails

Database Architecture Explained:

```

graph TD
    FS[Fraud Service] --> PG[(Primary PostgreSQL Writes)]
    PG --> SR[Streaming Replication]
    SR --> SB[(Standby DB)]

    AUDIT[Auditors] --> LB[Load Balancer]
    LB --> RR1[(Read Replica 1 Queries)]
    LB --> RR2[(Read Replica 2 Reports)]
    LB --> RC[(Redis Cache Hot data)]

    PG -. -> RR1
    PG -. -> RR2

    style FS fill:#e3f2fd,color:#000000
    style PG fill:#e8f5e8,color:#000000
    style SR fill:#f3e5f5,color:#000000
    style SB fill:#e8f5e8,color:#000000
    style AUDIT fill:#fff3e0,color:#000000
    style LB fill:#f9f9f9,color:#000000
    style RR1 fill:#e8f5e8,color:#000000

```

```
style RR2 fill:#e8f5e8,color:#000000
style RC fill:#ffebee,color:#000000
```

Scaling Strategy Breakdown:

1. Write Performance:

- **Primary Database:** Handles all fraud case creation (10K+ writes/sec)
- **Connection Pooling:** 200+ connections with pgBouncer
- **Batch Inserts:** Group related data for efficiency

2. Read Performance:

- **Read Replicas:** Dedicated servers for auditor queries (no impact on writes)
- **Redis Cache:** Hot data (recent cases, user sessions) in memory
- **Query Optimization:** Indexes on common search patterns

3. Data Management:

- **Partitioning:** Monthly partitions keep query performance high
- **Archive Strategy:** Move resolved cases >2 years to cold storage
- **Compression:** PostgreSQL native compression for historical data

4. Auditor Interface - Human-AI Collaboration

Why Human Auditors Are Essential: Even with 98% precision, a bank processing 1M transactions/day still gets **20,000 false positives daily**. Human investigators provide:

- **Context understanding:** "Customer just moved, large purchases are normal"
- **Complex pattern recognition:** Multi-step fraud schemes spanning weeks
- **Customer interaction:** Phone calls to verify suspicious activity
- **Regulatory compliance:** Human oversight required by banking regulations

How Auditors Work with the System:

```
graph LR
    ML[ML Model] --> FA[Fraud Alert]
    FA --> AD[Auditor Dashboard]
    AD --> INV[Investigation]
    INV --> RES[Resolution]

    ML -. -> AF[Auto-flags suspicious]
    FA -. -> PR[Prioritizes by risk level]
    AD -. -> GE[Gathers evidence]
    INV -. -> CU[Customer unblocked]





    style ML fill:#e3f2fd,color:#000000
```

```
style FA fill:#ffebee,color:#000000
style AD fill:#f3e5f5,color:#000000
style INV fill:#fff3e0,color:#000000
style RES fill:#e8f5e8,color:#000000
style AF fill:#f9f9f9,color:#000000
style PR fill:#f9f9f9,color:#000000
style GE fill:#f9f9f9,color:#000000
style CU fill:#f9f9f9,color:#000000
```

Dashboard Features Explained:

1. Real-time Case Dashboard:

Risk Level Summary:

Priority	Count	Icon
CRITICAL	47	
HIGH	156	
MEDIUM	312	
LOW	89	

Dashboard Filters:

Filter Type	Options	Current Selection
Status	All, Pending, Investigating, Resolved	All
Assignment	All cases, Assigned to me, Unassigned	Assigned to me ✓
Date Range	Last 24h, Last 7d, Last 30d, Custom	Last 24h
Amount Filter	Any, >\$1K, >\$10K, >\$50K, Custom	Custom: >\$____

Case List:

Case ID	Account	Amount	Type	Confidence	Age	Actions
FR_001234	ACC_567890	\$15,847.32	TRANSFER	0.94	2h	[Investigate]
FR_001235	ACC_234567	\$8,923.11	CASH_OUT	0.87	4h	[Assign to me]

2. Investigation Workflow:

```
graph TD
    P[PENDING] --> A[Assign]
    A --> I[INVESTIGATING]
    I --> AE[Add Evidence]
    AE --> R[RESOLVED]
```

```
I --> RM[Request more info]
RM --> CC[Contact customer]
CC --> ML[Mark as legitimate/fraud]
ML --> R
```

```
style P fill:#fff3e0,color:#000000
style A fill:#e3f2fd,color:#000000
style I fill:#f3e5f5,color:#000000
style AE fill:#e8f5e8,color:#000000
style R fill:#c8e6c9,color:#000000
style RM fill:#f9f9f9,color:#000000
style CC fill:#f9f9f9,color:#000000
style ML fill:#f9f9f9,color:#000000
```

3. Transaction Detail View:

- **Full Transaction Context:** Account history, recent patterns, device info
- **ML Explanation:** "Flagged because: balance inconsistency + unusual time"
- **Similar Cases:** "3 similar patterns resolved as fraud in past week"
- **Customer Profile:** VIP status, complaint history, typical behavior

4. Bulk Actions for Efficiency:

- **Pattern Recognition:** "Mark all TRANSFER transactions from IP 192.168.1.1 as fraud"
- **Account-based:** "Review all transactions from account ACC_123456"
- **Time-based:** "Investigate all high-risk cases from last 4 hours"

API Endpoints with Real-World Usage:

```
# Get pending cases assigned to current investigator
GET /cases?status=PENDING&assigned_to=me&sort=confidence_desc&page=1

# Get case details with full context
GET /cases/FR_001234
Response: {
  "case_id": "FR_001234",
  "transaction": {...},
  "ml_explanation": "Primary factors: balance_error_src=1.0,
unusual_time=0.8",
  "similar_cases": [...],
  "account_history": [...],
  "customer_profile": {...}
}

# Update case with investigation notes
PUT /cases/FR_001234
Body: {
  "status": "INVESTIGATING",
  "notes": "Called customer - confirmed legitimate. Large purchase for
home renovation.",
  "resolution": "FALSE_POSITIVE"
```

```
}

# Get analytics for management dashboard
GET /analytics/fraud-trends?period=last_30_days
Response: {
  "cases_created": 45234,
  "resolution_rate": 0.94,
  "false_positive_rate": 0.15,
  "avg_resolution_time_hours": 6.2,
  "top_fraud_patterns": [...]
}
```

Scalability & Reliability Considerations

Horizontal Scaling - Handling Traffic Spikes

Why Auto-Scaling is Critical: Banking traffic is **highly unpredictable**:

- **Black Friday:** 10x normal transaction volume
- **Payroll Days:** 5x volume at month-end
- **System Outages:** Catch-up processing creates massive backlogs

Auto-Scaling Strategy:

Normal Load	Peak Load
↓	↓
3 consumers	→ 12 consumers
2 API pods	→ 10 API pods

Scaling Triggers & Thresholds:

- **Kafka Lag > 10,000 messages:** Add 2 consumer instances
- **API Response Time > 100ms:** Add 2 API replicas
- **Database CPU > 80%:** Promote read replica to writer
- **Fraud Detection Queue > 1,000:** Emergency scaling protocol

Fault Tolerance - Zero Downtime Requirements

Why 99.99% Uptime is Required:

- **Revenue Impact:** Bank loses \$50,000/minute during payment system downtime
- **Regulatory Penalties:** Compliance violations for missed fraud detection
- **Customer Trust:** Payment failures drive customers to competitors

Multi-Layer Fault Tolerance:

1. Service Level (Kubernetes):

- **Auto-restart failed containers:** Liveness probe checks `/health` endpoint every 10 seconds
- **Health monitoring:** Readiness probe checks `/ready` endpoint every 5 seconds

- **High availability:** 3 replicas minimum to prevent single points of failure
- **Node distribution:** Pod anti-affinity ensures replicas spread across different nodes

2. Data Level:

- **Primary Database:** Handles all write operations with synchronous replication to standby
- **Standby Database:** Hot standby with automatic failover in 30 seconds if primary fails
- **Read Replicas:** 2 separate instances for query load distribution via async replication
- **Connection Management:** Automatic traffic routing to healthy database instances

3. Message Level (Kafka):

- **Replication Factor 3:** Each message stored on 3 brokers
- **In-Sync Replicas:** Minimum 2 replicas must acknowledge writes
- **Consumer Groups:** If one consumer fails, others take over its partitions

4. Model Failures - Graceful Degradation:

- **Primary ML Model:** XGBoost model provides main fraud detection capability
- **Fallback System:** Rule-based detection takes over when ML model fails
- **Automatic Alerting:** ML team immediately notified of model failures for rapid response
- **Service Continuity:** System continues fraud detection with reduced accuracy rather than failing completely

Security - Banking-Grade Protection

Multi-Layer Security Model:

- **Internet Layer:** DDoS protection shields against volumetric attacks
- **Web Application Firewall (WAF):** Filters malicious requests and application-layer attacks
- **Load Balancer:** Implements rate limiting to prevent abuse and overload
- **API Gateway:** Handles authentication/authorization using JWT tokens with role-based permissions
- **Service Layer:** Operates within secure internal network with restricted communication
- **Database Layer:** All data encrypted at rest using AES-256 encryption

Security Implementation Details:

1. Network Security:

- **VPC Isolation:** All services in private subnets (no internet access)
- **Security Groups:** Strict port/protocol restrictions
- **Network ACLs:** Additional layer of subnet-level filtering
- **VPN/Transit Gateway:** Secure connection to bank's core systems

2. Authentication & Authorization:

- **JWT Tokens:** Stateless authentication with 8-hour expiration for security
- **Role-Based Access:** Users assigned specific roles (fraud_analyst, manager, admin)
- **Granular Permissions:** Fine-grained permissions for viewing, updating, and contacting customers
- **Session Management:** Automatic token refresh and secure logout procedures

3. Data Protection:

- **Encryption in Transit:** TLS 1.3 for all communication
- **Encryption at Rest:** AES-256 for database, S3, message queues
- **PII Tokenization:** Credit card numbers replaced with tokens
- **Data Masking:** Sensitive fields masked in logs/analytics

4. Compliance & Auditing:

- **SOX Compliance:** All financial data access logged
- **GDPR Compliance:** Customer data handling and deletion rights
- **PCI DSS:** Payment card data security standards
- **Audit Trails:** Immutable log of all system actions

Monitoring & Alerts - Proactive Issue Detection

Why Comprehensive Monitoring is Essential: Fraud systems are **mission-critical** - issues must be detected and resolved within minutes:

3-Tier Alert Strategy:

Tier 1 (Immediate - < 5 min response)

- └ System down/Model failures → Page on-call engineer
- └ High fraud detection rate → Alert fraud team lead
- └ Database connection failures → Auto-remediate + alert

Tier 2 (Urgent - < 30 min response)

- └ High API latency → Alert platform team
- └ Unusual fraud patterns → Alert fraud analysts
- └ Kafka lag growing → Alert data engineering

Tier 3 (Important - < 2 hour response)

- └ Disk space trending high → Alert infrastructure
- └ Model accuracy declining → Alert ML team
- └ Customer complaints trending up → Alert business team

Monitoring Dashboard (Real-time):



BUSINESS METRICS

- └ Fraud Detection Rate: 2.3%
- └ False Positive Rate: 1.8%
- └ Cases Resolved: 1,247/day
- └ Customer Impact: 12 blocked



FRAUD TRENDS

- └ Transfer fraud ↑ 15%
- └ Geographic: Nigeria ↑ 30%
- └ Peak hours: 2-4 PM EST



SYSTEM HEALTH

- └ API Latency: 45ms (p95)
- └ Kafka Lag: 234 messages
- └ Database Connections: 87/200
- └ Error Rate: 0.02%



ACTIVE ALERTS

- └ High-risk case spike
- └ Investigator queue full

Integration with Business Systems:

- **Slack Integration:** Real-time alerts to #fraud-ops channel
- **Email Escalation:** Manager alerts for critical issues
- **Dashboard Screens:** Wall-mounted displays in fraud operations center
- **Mobile Apps:** Push notifications for on-call investigators

This architecture balances simplicity with production readiness, ensuring reliable real-time fraud detection while providing auditors with efficient investigation tools.